

Spotify Data Visualization

Elad Golan

2022-05-23

```
# install.packages(c("tidyverse", "ggmosaic", "glue", "tidymodels", "glmnet", "randomForest", "kernlab")
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr 0.3.4
## v tibble 3.1.6       v dplyr 1.0.8
## v tidyr 1.2.0        v stringr 1.4.0
## v readr 2.1.2        v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(ggmosaic)
library(glue)
library(tidymodels)
```

```
## Registered S3 method overwritten by 'tune':
```

```
##   method                from
##   required_pkgs.model_spec parsnip
```

```
## -- Attaching packages ----- tidymodels 0.1.4 --
```

```
## v broom      0.7.12      v rsample      0.1.1
## v dials      0.1.0       v tune         0.1.6
## v infer      1.0.0       v workflows    0.2.4
## v modeldata  0.1.1       v workflowsets 0.1.0
## v parsnip    0.2.0       v yardstick    0.0.9
## v recipes    0.2.0
```

```
## -- Conflicts ----- tidymodels_conflicts() --
```

```
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed()  masks stringr::fixed()
## x dplyr::lag()       masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step()   masks stats::step()
## x tune::tune()      masks parsnip::tune()
## * Dig deeper into tidy modeling with R at https://www.tmr.org
```

```
library(kernlab)
```

```
##
```

```
## Attaching package: 'kernlab'
```

```
## The following object is masked from 'package:scales':
```

```

##
##      alpha
## The following object is masked from 'package:purrr':
##
##      cross
## The following object is masked from 'package:ggplot2':
##
##      alpha
library(ggradar)
library(ragg)
library(showtext)

## Loading required package: sysfonts
## Loading required package: showtextdb
library(tvthemes)
library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
## The following objects are masked from 'package:yardstick':
##
##      precision, recall, sensitivity, specificity
## The following object is masked from 'package:purrr':
##
##      lift
library(ggwordcloud)
library(tm)

## Loading required package: NLP
##
## Attaching package: 'NLP'
## The following object is masked from 'package:ggplot2':
##
##      annotate
library(stringi)

https://github.com/rfordatascience/tidytuesday/tree/master/data/2020/2020-01-21
spotify_songs <- read_csv(https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2020/2020-01-21/spotify\_songs.csv)

## Rows: 32833 Columns: 23
## -- Column specification -----
## Delimiter: ","
## chr (10): track_id, track_name, track_artist, track_album_id, track_album_na...
## dbl (13): track_popularity, danceability, energy, key, loudness, mode, spec...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

```
spotify_songs %>% count(playlist_genre)
```

```
## # A tibble: 6 x 2
##   playlist_genre      n
##   <chr>           <int>
## 1 edm             6043
## 2 latin           5155
## 3 pop             5507
## 4 r&b             5431
## 5 rap             5746
## 6 rock            4951
```

Attention: Apparently each song can repeat a few times with a few genres!

```
spotify_songs %>% count(track_id, sort = TRUE)
```

```
## # A tibble: 28,356 x 2
##   track_id          n
##   <chr>           <int>
## 1 7BKLCZ1jbUBVqRi2FV1TVw    10
## 2 14s0S5L36385FJ30L8hew4     9
## 3 3eekarcy7kvN4yt5ZFzltW     9
## 4 0nbXyq5TXYPc07pr3N8S4I     8
## 5 0qaWEvPkts34WF68r8Dzx9     8
## 6 0rIAC4PXANcKmitJfoqmVm     8
## 7 0sf12qNH5qcw8qpgymF0qD     8
## 8 2b8f0ow8UzyDFAE27Yh0ZM     8
## 9 2Fxmhks0bxGSBdJ92vM42m     8
## 10 2tnVG71enUj33Ic2nFN6kZ     8
## # ... with 28,346 more rows
```

You can either ignore this for now or sample a genre for each song, e.g. with:

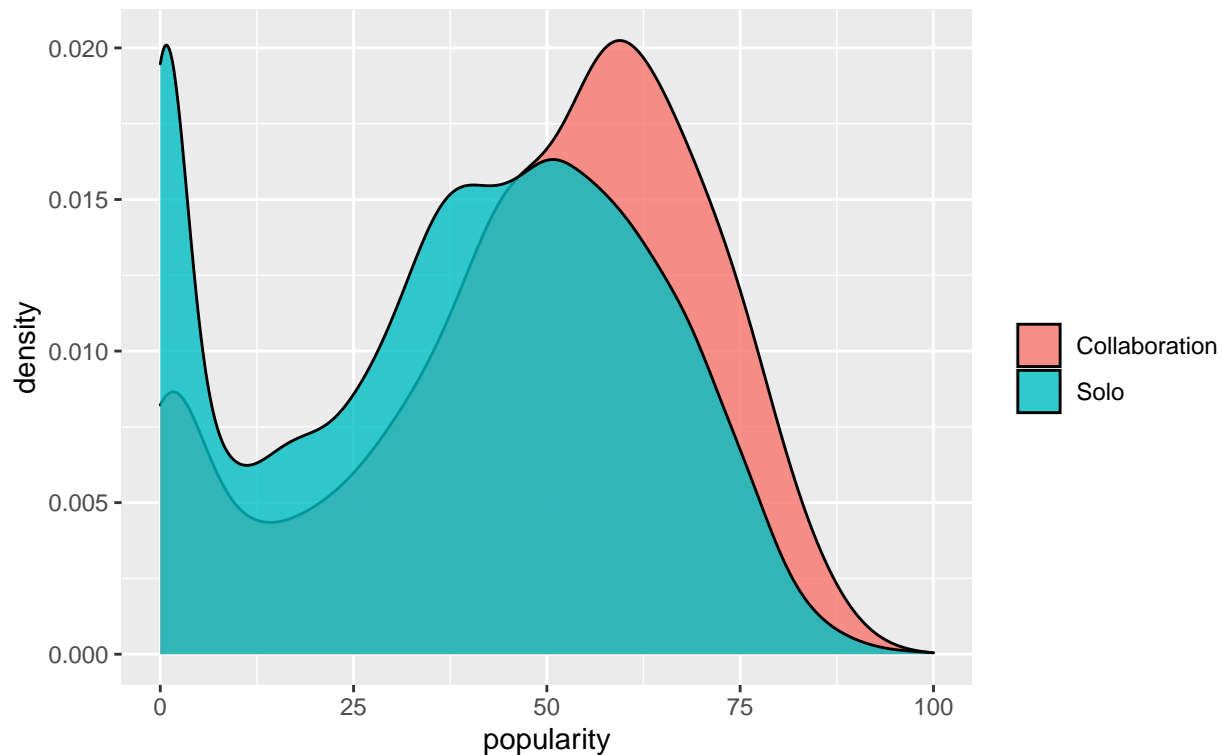
```
spotify_songs <- spotify_songs %>%
  group_by(track_id) %>%
  slice_sample(n = 1) %>%
  ungroup()
```

```
p4 <- ggplot(spotify_songs %>% select(track_name, track_album_name, track_popularity) %>% filter(!(grepl
```

```
p4 + geom_density(aes(fill=feat), alpha=0.8) + labs(title="Density plot",
  subtitle="Does collaboration between artists increase the chance that the song will be popular",
  x="popularity",
  fill="")
```


Density plot

Does collaboration between artists increase the chance that the song will be popular?



```
# I cut the years into 4-year groups, except for 85-88 which represents a 3-year period.
ggplot(spotify_songs %>% select(playlist_genre, track_album_release_date, track_popularity) %>%
  mutate(track_album_release_year = cut(as.numeric(substr(track_album_release_date,1,4)), breaks = 16))
  group_by(playlist_genre, track_album_release_year) %>%
  summarise(mean_popularity = mean(track_popularity), n = n(), .groups = 'drop') %>% filter(n >= 10),
  aes(track_album_release_year, mean_popularity, size = n, colour = playlist_genre)) + geom_point(alpha
    scale_x_discrete(labels = c("61-65", "65-69", "69-73", "73-77", "77-81", "81-85", "85-88", "88-92", "92-
      "96-00", "00-04", "04-08", "08-12", "12-16", "16-20")) +
  labs(title = "songs` popularity over time by genre", y = "popularity", x = "years") + theme_avatar() +
  theme(text = element_text(family = "serif"), plot.title = element_text(size = 14, face = "bold", hjust = 0),
    legend.position = "bottom")
```


```
## Scale for 'colour' is already present. Adding another scale for 'colour',
## which will replace the existing scale.
```



```
font_add_google("pacifico")
showtext_auto()

ggradar(spotify_songs %>% group_by(playlist_genre) %>%
  summarise(across(c(danceability, instrumentalness, speechiness, energy, valence), mean))
  , font.radar = "pacifico", group.point.size = 3, axis.label.size = 4, grid.label.size = 3,
  values.radar = c("0", "0.425", "0.85"), grid.min = 0, grid.mid = 0.4, grid.max = 0.85,
  legend.text.size = 10) + labs(title = "Radar plot of songs genre") +
  theme(legend.background = element_blank(), legend.key = element_rect(fill = NA, color = NA),
  plot.background = element_rect(fill = "#fbf9f4", color = "#fbf9f4"),
  panel.background = element_rect(fill = "#fbf9f4", color = "#fbf9f4"),
  plot.title.position = "plot",
  text = element_text(family = "pacifico"),
  plot.title = element_text(
    size = 15,
    face = "bold",
    color = "#2a475e"
  )
  )
)
```

Radar plot of songs genre



#filter the artists' name only for those who their first album release after 2005.

```
names_filter <- spotify_songs %>% mutate(as.numeric(substr(track_album_release_date,1,4))) %>%
  filter(track_album_release_date <= 2005) %>% pull(unique(track_artist))
```

#build the plot based the name_filter

```
p <- ggplot(spotify_songs %>% mutate(type= case_when(grepl(c("Remix"),track_name) |
  grepl(c("Remix"), track_album_name) ~ "Remix",
  TRUE ~ "Original"))) %>% group_by(track_artist, type) %>%
  summarise(total_type = n(), .group = 'drop') %>% inner_join(spotify_songs %>% group_by(track_artist,
  summarise(total_songs = n()), by = "track_artist") %>%
  select(track_artist, type, total_type, total_songs) %>%
  mutate(ratio = case_when(type == "Remix" ~ paste0(round(total_type/total_songs, 2)*100, "%"), TRUE ~
  filter(total_songs > 30, !track_artist %in% names_filter),
  aes(x = track_artist, y = total_type ,fill = type))
```

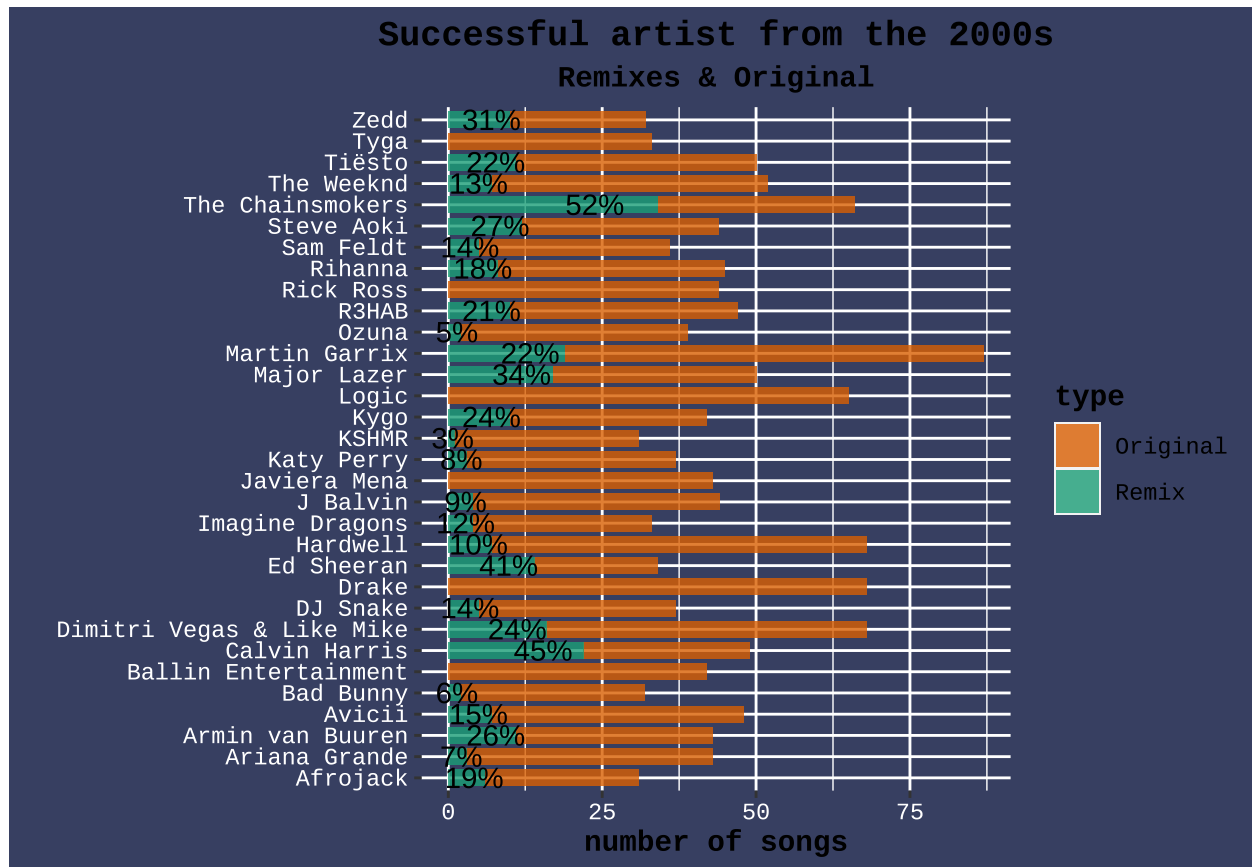
`summarise()` has grouped output by 'track_artist'. You can override using the
`.groups` argument.

```
p + geom_col(alpha=0.8 , width=.8) + coord_flip() +
  geom_text(aes(label=ratio), position =position_stack(vjust = .7)) +
  scale_fill_brewer(palette = "Dark2", direction = -1) + labs( title = "Successful artist from the 2000s",
  subtitle = "Remixes & Original" ,
  x = "", y = "number of songs") +
  theme(text = element_text(family = "mono"), plot.title = element_text(face = "bold", hjust = 0.5),
  title = element_text(face = "bold"), plot.subtitle = element_text(hjust = 0.5) ,
  plot.background = element_rect(fill = "#373F61", color = "#373F61"),
```

```

panel.background = element_rect(fill = "#373F61", color = "#373F61"),
legend.background = element_rect(fill = "#373F61", color = "#373F61"),
axis.text.x = element_text(color = "white"), axis.text.y = element_text(color = "white"))

```



```

#function for cleaning the songs' names from punctuation, stop words, digits and irrelevant words as Re
clean_fun <- function(name){
  name <- removeWords(str_replace_all_regex(str_to_lower(name) ,
                                             c("\\(feat\\.\\.*", "\\feat\\.\\.*", "\\[feat\\.\\.*", "\\-\\.\\.*", "\\(\\.\\.*)", "",
                                             vectorize=FALSE), stopwords()))

  name <- str_remove_all(name, "[[:punct:]]")
  return(str_remove_all(name, "[[:digit:]]"))
}

#applying the clean_fun on pop songs.
pop_songs_clean <- spotify_songs %>% filter(playlist_genre == "pop") %>%
  select(track_name) %>%
  mutate(track_name = clean_fun(track_name)) %>% pull()

#making a DF of the freq of the words in the pop songs' names.
counter_words <- as.data.frame(table(unlist(strsplit(pop_songs_clean, " "))))

set.seed(42) # The way the words appear in the chart is random and we would like to get identical resul

#I filtered for words that appear more than 5 times so that the chart does not become too cluttered.
counter_words %>% filter(Freq > 5) %>%

```

